

The work you submit in this course is required to be the result of your individual effort **only**. You may discuss concepts and ideas with others, but **you must program individually**. You should **never observe another student's code**, from this or previous semesters.

Your code will be a simulation because it will run in a normal process. We provide two files, `os.c` and `os.h`, which contain helper functions that simulate some OS functionality that your code will need to call. For your convenience, there's also a `main()` function demonstrating usage of the code. However, the provided `main()` only exercises your code in trivial ways. We recommend that you change it to thoroughly test the functions that you implemented.


Our simulated OS targets an imaginary 64-bit x86-like CPU. When talking about addresses (virtual or physical), we refer to the least significant bit as bit 0 and to the most significant bit as bit 63.

**Virtual addresses** The virtual address size of our hardware is 64 bits, of which only the lower 57 bits are used for translation. The top 7 bits are guaranteed to be identical to bit 56, i.e., they are either all ones or all zeroes. The following depicts the virtual address layout:

[illegible]

**Physical addresses** The physical address size of our hardware is also 64 bits.

**Page table structure** The page/frame size is 4KB (4096 bytes). Page table nodes occupy a physical page frame, i.e., they are 4KB in size. The size of a page table entry is 64 bits. Bit 0 is the valid bit. Bits 1–11 are unused and must be set to zero. (This means that our target CPU does not implement page access rights.) The top 52 bits contain the page frame number that this entry points to. The following depicts the PTE format:



**Number of page table levels** To successfully complete the assignment, you must answer to yourself: how many levels are there in our target machine's multi-level page table? As mentioned, assume that only the lowest 57 bits of the virtual address are used for translation.

## 1.2 OS physical memory manager

To write code that manipulates page tables, you need to be able to perform the following: (1) obtain the page number of an unused physical page, which marks it as used; (2) return a physical page that you don't need back to the OS, which marks it back as unused; and (3) obtain the kernel virtual address of a given physical address. The provided `os.c` contains functions simulating this functionality:

1. Use the following function to allocate a physical page (also called *page frame*):

```
uint64_t alloc_page_frame(void);
```

This function returns the **physical page number** of the allocated page. In this assignment, you do not need to free physical pages. If `alloc_page_frame()` is unable to allocate a physical page, it will exit the program. The content of the allocated page frame is all zeroes.

2. Use the following function to free a page frame. This function takes the **physical page number** of the freed page:

```
void free_page_frame(uint64_t ppn);
```

The argument was previously returned from `alloc_page_frame` and was not yet freed. Otherwise the behavior is undefined.

3. Use the following function to obtain a pointer (i.e., virtual address) to a **physical address**:

```
void* phys_to_virt(uint64_t phys_addr);
```

The valid inputs to `phys_to_virt()` are addresses that reside in physical pages that were previously returned by `alloc_page_frame()` and not yet freed using `free_page_frame()`. If it is called with an invalid input, it returns NULL.

## 2 Assignment description

Implement the following two functions in a file named `pt.c`. This file should `#include "os.h"` to obtain the function prototypes.

1. A function to create/destroy virtual memory mappings in a page table:

```
void page_table_update(uint64_t pt, uint64_t vpn, uint64_t ppn);
```

This function takes the following arguments:

- (a) **pt**: The **physical page number** of the page table root (this is the physical page that the page table base register in the CPU state will point to). You can assume that **pt** has been previously returned by `alloc_page_frame()`.
- (b) **vpn**: The **virtual page number** the caller wishes to map/unmap.
- (c) **ppn**: Can be one of two cases. If **ppn** is equal to a special `NO_MAPPING` value (defined in `os.h`), then **vpn**'s mapping (if it exists) should be destroyed. Otherwise, **ppn** specifies the **physical page number** that **vpn** should be mapped to. If after the change, all the mappings are `NO_MAPPING`, you should free the page table node. If the node is freed, then it should be marked with `NO_MAPPING` in its parent. Notice, the root node should never be freed.

2. A function to query the mapping of a **virtual page number** in a page table:

```
uint64_t page_table_query(uint64_t pt, uint64_t vpn);
```

This function returns the **physical page number** that **vpn** is mapped to, or `NO_MAPPING` if no mapping exists. The meaning of the **pt** argument is the same as with `page_table_update()`.

You can implement helper functions for your code, but make sure to implement them in your `pt.c` file. You may not submit additional files (not even header files).

**IMPORTANT:** A page table node should be treated as an array of `uint64_ts`.

## 2.1 Something to think about (no need to submit)

What is the cost of freeing a page table node once all the PTEs it contains become invalid? How efficient is your approach (i.e., how much overhead does it add to each page table update)? If we wouldn't have freed the tables, what would be the cost? How can we avoid wasting frames for a process that allocated a lot and then freed its mapping if we do not implement deletion?

## 3 Submission instructions

1. Submit just your `pt.c` file. (We will test it with our own `main` function.)
2. The program must compile cleanly (no errors or warnings) when the following command is run in a directory containing the source code files:

```
gcc -O3 -Wall -std=c11 os.c pt.c
```